

ELEC2870 - Machine learning: regression and dimensionality reduction

Nonlinear regression with Multi-Layer Perceptrons

Michel Verleysen

Machine Learning Group

Université catholique de Louvain

Louvain-la-Neuve, Belgium

michel.verleysen@uclouvain.be


Outline

- Motivation
- Single-layer nonlinear regression
- Multi-layer perceptron
 - Model
 - MLP with threshold units
 - Number of layers
 - Learning
 - Error back-propagation
 - Weight adjustment
- Applications

Nonlinear regression: motivation

- Remember the linear model

$$y = \mathbf{w}^T \mathbf{x}$$





- and the sum-of-squares criterion

$$E = \frac{1}{P} \sum_{p=1}^P (t^p - y^p)^2 = \frac{1}{P} \sum_{p=1}^P (t^p - \mathbf{w}^T \mathbf{x}^p)^2$$

- The influence of a *large* single error on this criterion is *very large* (because the error is squared)

Nonlinear regression: motivation

- When the dataset is quite small, a single outlier may have a dramatic influence:




- This is true both in classification and regression!


Nonlinear regression: motivation

- Remember also that a linear model simply does not the expected job...

– Linear model



– Nonlinear model




Outline

- Motivation
- Single-layer nonlinear regression
- Multi-layer perceptron
 - Model
 - MLP with threshold units
 - Number of layers
 - Learning
 - Error back-propagation
 - Weight adjustment
- Applications

Single-layer nonlinear regression

- Model identical to linear regression, but with *nonlinear activation function*

$$y = \sigma(\mathbf{w}^T \mathbf{x})$$



- The error function becomes

$$E = \frac{1}{P} \sum_{k=1}^P \left(t^k - \sigma(\mathbf{w}^T \mathbf{x}^k) \right)^2$$

- And the stochastic gradient descent rule


$$\mathbf{w}(t+1) = \mathbf{w}(t) + \frac{2}{P} \alpha \left(t^k - \sigma(\mathbf{w}(t)^T \mathbf{x}^k) \right) \mathbf{x}^k \frac{\partial \sigma}{\partial p} \bigg|_{p=\mathbf{w}(t)^T \mathbf{x}^k}$$

(sometimes called the *generalized delta rule*)


Nonlinear activation functions

- Commonly used non-linear activation function

hyperbolic tangent




logistic function




- any slope
- slope can be a parameter
- logistic function: used to estimate posterior probabilities

Nonlinear effect on outliers

without non-linearity



with non-linearity




- Why? Because $|\sigma(\mathbf{w}^T \mathbf{x}^k)|$ never exceeds one, therefore $t^k - \sigma(\mathbf{w}^T \mathbf{x}^k)$ is almost 0 for well-classified points
- Not so simple in regression

Outline

- Motivation
- Single-layer nonlinear regression
- Multi-layer perceptron
 - Model
 - MLP with threshold units
 - Number of layers
 - Learning
 - Error back-propagation
 - Weight adjustment
- Applications

Multi-Layer Perceptron (MLP)


- several layers of weights and activation units



Multi-Layer Perceptron

$$y_k(\mathbf{x}) = h \left(\sum_{i=0}^M w_{ki}^{(2)} g \left(\sum_{j=0}^D w_{ij}^{(1)} x_j \right) \right)$$

$$\mathbf{y}(\mathbf{x}) = \mathbf{h}(\mathbf{w}^{(2)} \mathbf{g}(\mathbf{w}^{(1)} \mathbf{x}))$$



- Convention: 2 layers of weights (in literature: sometimes 3 layers of *units* or *neurons*)
- g and h can be threshold (sign) units or continuous ones
- h can be linear but not g (otherwise only one layer)

Multi-Layer Perceptron

- How many layers
 - can we use ?
 - should we use ?
- In theory:
 - Any number of layers (see back-propagation algorithm)
- In practice:
 - More layers means more $w_{ij}^{(l)}$ parameters
 - Is it really needed?
 - How many layers do we need to approximate any function?

Outline

- Motivation
- Single-layer nonlinear regression
- Multi-layer perceptron
 - Model
 - MLP with threshold units
 - Number of layers
 - Learning
 - Error back-propagation
 - Weight adjustment
- Applications

MLP with threshold units



Warning: this is NOT the traditional MLP !

- We discuss the *MLP with threshold units* first to get an intuition about the number of layers
- But then we forget it immediately!
- *Why? Because threshold units means discontinuities, so no algorithm...*
- Outputs:
 - threshold units (all layers) → binary outputs
- Inputs:
 - binary inputs: Boolean function network
 - 2 layers (max.) for any function
 - look-up table without generalisation
 - continuous inputs (classification)
 - data become binary after 1st layer
 - shape of decision boundaries depend on # layers



MLP with threshold units

Warning: this is NOT the traditional MLP !



linear



convex




any



Multi-Layer Perceptron

- General case (no more threshold units!)
- non-linear activation function: sigmoid or hyperbolic tangent
 - at least for “hidden” layers
 - output layer can be linear (otherwise limited output range)
- used for
 - approximation of functions
 - classification







Outline

- Motivation
- Single-layer nonlinear regression
- Multi-layer perceptron
 - Model
 - MLP with threshold units
 - Number of layers
 - Learning
 - Error back-propagation
 - Weight adjustment
- Applications

Number of layers

- A 3-layers network can approximate
 - any function
 - with any precision
- Indeed:
 - A 2-layers network can approximate a local function



2 layers:

(a) sigmoid

(b) sum of 2 sigmoids

(c) sum of 4 sigmoids

(d) sigmoid of sum
(bell-shaped local
function)

- A 3-layers network can thus approximate any sum of local functions

Number of layers

- That was intuition...
- Mathematically, it can be proven that

A **2**-layer MLP can approximate arbitrarily well any (functional) continuous mapping, provided the number M of hidden units is sufficiently large

- This is called the *universal approximation property*
- It is theory: in practice, if M is too high, overfitting!
- It also valid for decision boundaries (classification)

Outline


- Motivation
- Single-layer nonlinear regression
- Multi-layer perceptron
 - Model
 - MLP with threshold units
 - Number of layers
 - Learning
 - Error back-propagation
 - Weight adjustment
- Applications

Learning in MLP

- Learning =
 - definition of an error criterion E
 - evaluation of derivatives of E w.r.t. parameters w
 - adjustments of parameters w according to derivatives

- Error criterion
$$E = \frac{1}{P} \sum_{p=1}^P E^p(y_1, \dots, y_C)$$

- Batch / stochastic learning
 - Batch learning (all samples together): E
 - on-line, stochastic (one sample at each iteration): E^p
- In the following, we omit p (stochastic learning)




$$y_k(\mathbf{x}) = h \left(\sum_{i=0}^M w_{ki}^{(2)} g \left(\sum_{j=0}^D w_{ij}^{(1)} x_j \right) \right)$$

Outline


- Motivation
- Single-layer nonlinear regression
- Multi-layer perceptron
 - Model
 - MLP with threshold units
 - Number of layers
 - Learning
 - Error back-propagation
 - Weight adjustment
- Applications

Error back-propagation

- Some notations




$$a_i^{(l)} = \sum_j w_{ij}^{(l)} z_j^{(l-1)}$$



Error back-propagation

$$a_i^{(l)} = \sum_j w_{ij}^{(l)} z_j^{(l-1)}$$

- Gradient descent: evaluation of $\frac{\partial E}{\partial w_{ij}^{(l)}}$



$$\frac{\partial E}{\partial w_{ij}^{(l)}} = \frac{\partial E}{\partial a_i^{(l)}} \frac{\partial a_i^{(l)}}{\partial w_{ij}^{(l)}}$$

$$= \delta_i^{(l)} z_j^{(l-1)}$$

$$\delta_i^{(l)} \equiv \frac{\partial E}{\partial a_i^{(l)}} \quad \text{to evaluate}$$

Error back-propagation

$$\delta_i^{(l)} \equiv \frac{\partial E}{\partial a_i^{(l)}} \quad \text{to evaluate}$$

- For output units


$$\delta_i^{(l)} = \frac{\partial E}{\partial a_i^{(l)}} = \frac{\partial E}{\partial y_i^{(l)}} \frac{\partial y_i^{(l)}}{\partial a_i^{(l)}}$$

$$= \frac{\partial E}{\partial y_i^{(l)}} g'(a_i^{(l)})$$

known

derivative of error criterion (known): $E = (t_i - y_i)^2$

$$\Rightarrow \frac{\partial E}{\partial y_i} = -2(t_i - y_i)$$




Error back-propagation

$$\delta_i^{(l)} \equiv \frac{\partial E}{\partial a_i^{(l)}} \quad \text{to evaluate}$$

- For hidden units

$$\begin{aligned} \delta_i^{(l)} &= \frac{\partial E}{\partial a_i^{(l)}} = \sum_k \frac{\partial E}{\partial a_k^{(l+1)}} \frac{\partial a_k^{(l+1)}}{\partial a_i^{(l)}} \\ &= \sum_k \left(\delta_k^{(l+1)} \frac{\partial \left(\sum_{j \in (l)} w_{kj}^{(l+1)} z_j^{(l)} \right)}{\partial a_i^{(l)}} \right) \\ &= \sum_k \left(\delta_k^{(l+1)} w_{ki}^{(l+1)} g'(a_i^{(l)}) \right) = g'(a_i^{(l)}) \sum_k \left(\delta_k^{(l+1)} w_{ki}^{(l+1)} \right) \end{aligned}$$



The error term (δ) is expressed as a combination of errors in the next layer

Error back-propagation

- Algorithm:

- Apply an input vector \mathbf{x}^k and propagate it through the network to evaluate all activations $z_i^{(l)}$ and neuron outputs $a_i^{(l)}$
- Evaluate error terms $\delta_i^{(o)}$ in output layer
- Back-propagate error terms $\delta_i^{(l)}$ to find error terms $\delta_i^{(l-1)}$
- Evaluate all derivatives $\frac{\partial E}{\partial w_{ij}^{(l)}} = \delta_i^{(l)} z_j^{(l-1)}$
- Adjust weights according to derivatives and a gradient descent scheme

Outline

- Motivation
- Single-layer nonlinear regression
- Multi-layer perceptron
 - Model
 - MLP with threshold units
 - Number of layers
 - Learning
 - Error back-propagation
 - Weight adjustment
- Applications

Weight adjustment

- Gradient descent is nice, but...
 - it can be very slow (when there are many parameters)
 - it is easily stuck in local minima
 - it is quite easy to do better...
- Now we have the derivatives $\frac{\partial E}{\partial w_{ij}^{(l)}}$, how to adjust weights?
- Notations:

– weight update $\delta \mathbf{w} \equiv \mathbf{w}(t+1) - \mathbf{w}(t)$

– gradient $\left(\frac{\partial E}{\partial \mathbf{w}} \right)^T \equiv \left(\frac{\partial E}{\partial w_1} \frac{\partial E}{\partial w_2} \dots \frac{\partial E}{\partial w_D} \right)$

– Hessian

$$H \equiv \left(\frac{\partial^2 E}{\partial \mathbf{w}^2} \right) \equiv \begin{pmatrix} \frac{\partial^2 E}{\partial w_1^2} & \frac{\partial^2 E}{\partial w_1 \partial w_2} & \dots & \frac{\partial^2 E}{\partial w_1 \partial w_D} \\ \vdots & \vdots & & \vdots \\ \frac{\partial^2 E}{\partial w_D \partial w_1} & \frac{\partial^2 E}{\partial w_D \partial w_2} & \dots & \frac{\partial^2 E}{\partial w_D^2} \end{pmatrix}$$

Indices are changed for simplicity of notations:

$$w_{ij}^{(l)}, \forall i, j, l$$

↓

$$w_i, 1 \leq i \leq D$$

Weight adjustment

- First-order methods :

$$\begin{aligned}
 E(\mathbf{w}(t+1)) &= E(\mathbf{w}(t)) + \left(\frac{\partial E}{\partial \mathbf{w}} \bigg|_{\mathbf{w}(t)} \right)^T (\mathbf{w}(t+1) - \mathbf{w}(t)) \\
 &= E(\mathbf{w}(t)) + \left(\frac{\partial E}{\partial \mathbf{w}} \bigg|_{\mathbf{w}(t)} \right)^T \delta \mathbf{w}
 \end{aligned}$$

- For a $\delta \mathbf{w}$ of a given magnitude, largest $|\delta E|$ is found when gradient and weight update vectors are parallel


- Adaptation rule: $\mathbf{w}(t+1) = \mathbf{w}(t) - \alpha \frac{\partial E}{\partial \mathbf{w}} \bigg|_{\mathbf{w}(t)}$

- This is gradient descent; nice, isn't it?

Weight adjustment

- First-order methods

- Note that $-\frac{\partial E}{\partial \mathbf{w}}$ does not point towards the minimum of E !



Weight adjustment

- First-order methods: improvements
 - Momentum : to avoid sharp changes in gradient direction (stochastic scheme, outliers, etc.)

$$\mathbf{w}(t+1) = \mathbf{w}(t) - \alpha \frac{\partial E}{\partial \mathbf{w}} \bigg|_{\mathbf{w}(t)} + \beta (\mathbf{w}(t) - \mathbf{w}(t-1))$$

$$\beta \approx 0.9$$

Weight adjustment

- First-order methods: improvements

- adaptive learning rate α

$$w_{ij}(t+1) = w_{ij}(t) - \alpha_{ij}(t) \left. \frac{\partial E}{\partial w_{ij}} \right|_{\mathbf{w}(t)}$$

$$\delta w_{ij}(t-1) \delta w_{ij}(t) > 0 \Rightarrow \alpha_{ij}(t+1) = \alpha_{ij}(t) + \kappa$$

$$\delta w_{ij}(t-1) \delta w_{ij}(t) < 0 \Rightarrow \alpha_{ij}(t+1) = \alpha_{ij}(t) - \kappa$$

- Need for maximum safe step size
- This is more or less the « delta-bar-delta » rule

Weight adjustment

- Second-order methods (Newton's method)

$$E(\mathbf{w}(t+1)) = E(\mathbf{w}(t)) + \left(\frac{\partial E}{\partial \mathbf{w}} \bigg|_{\mathbf{w}(t)} \right)^T \delta \mathbf{w} + \frac{1}{2} \delta \mathbf{w}^T \frac{\partial^2 E}{\partial \mathbf{w}^2} \bigg|_{\mathbf{w}(t)} \delta \mathbf{w}$$

- Stationary point (zero derivative) of this quadratic form:

$$\delta \mathbf{w} = - \left(\frac{\partial^2 E}{\partial \mathbf{w}^2} \bigg|_{\mathbf{w}(t)} \right)^{-1} \frac{\partial E}{\partial \mathbf{w}} \bigg|_{\mathbf{w}(t)}$$

- Adaptation rule:
 - minimum: move in this direction
 - maximum or saddle point: move in conjugate direction
 - Many "conjugate gradient" rules (some are line searches)

Weight adjustment

- Second-order methods (Newton's method)

$$\delta \mathbf{w} = - \left(\frac{\partial^2 E}{\partial \mathbf{w}^2} \Big|_{\mathbf{w}(t)} \right)^{-1} \frac{\partial E}{\partial \mathbf{w}} \Big|_{\mathbf{w}(t)}$$


- Hessian matrix difficult to compute → various approximations
 - Successive estimations
 - Diagonal terms only (quasi-Newton)
- Levenberg–Marquardt is another efficient algorithm

Weight adjustment


- What about the size of the step?
- Stochastic way: small α
 - but if too small: slow
 - and if too large: jumps over the minimum
- Other solution: line search $\mathbf{w}(t+1) = \mathbf{w}(t) + \lambda(t)\mathbf{d}(t)$
 - go as far as possible in the chosen direction $\mathbf{d}(t)$
 - Implicit use of the Hessian matrix
 - If $d(t)$ is the gradient at each time step: $\mathbf{d}(t) = \left. \frac{\partial E}{\partial \mathbf{w}} \right|_{\mathbf{w}(t)}$
$$\frac{\partial}{\partial \lambda} E(\mathbf{w}(t) + \lambda \mathbf{d}(t)) = 0 \quad \text{thus} \quad \left. \frac{\partial E}{\partial \mathbf{w}} \right|_{\mathbf{w}(t+1)} \mathbf{d}(t) = 0$$
 - So successive directions are always orthogonal!

Weight adjustment

- Successive directions are orthogonal (line search)
 - Not good for speed of convergence



- Use *conjugate* gradients:
New direction is chosen so that the component of the gradient $\frac{\partial E}{\partial \mathbf{w}}$ parallel to the previous direction remains 0
 - Conjugate gradients make implicit use of the Hessian matrix



Outline

- Motivation
- Single-layer nonlinear regression
- Multi-layer perceptron
 - Model
 - MLP with threshold units
 - Number of layers
 - Learning
 - Error back-propagation
 - Weight adjustment
- Applications

Applications

- Some (old...) applications of MLP
- In fact some applications of nonlinear regression with machine learning! (not specific to MLP)
 1. A standard nonlinear model
 2. Several ways to use similar models for the same goal
 3. The interpolation-extrapolation question

Application: function approximation


- Standard multivariate regression (function approximation)
- Application to process optimisation (estimation of physical properties, based on chemical composition and process parameters)

$$\alpha = f(C, Si, Mn, P, S, Al, N, Cu, Cr, Ni, Sn, V, Mo, Ti, Nb, B, d, b, Ti, Tf)$$

- “relative yield stress” α of different steel qualities
- inputs:
 - 16 chemical additives
 - plate’s thickness d and width b
 - temperatures before (T_i) and after (T_f) rolling

From “Estimating material properties for process optimization”, T. Poppe & T. Martinetz (Siemens AG), in ICANN’93 (Amsterdam, The Netherlands) proceedings, S. Gielen & B. Kappen eds., Springer-Verlag, 1993

Application: function approximation



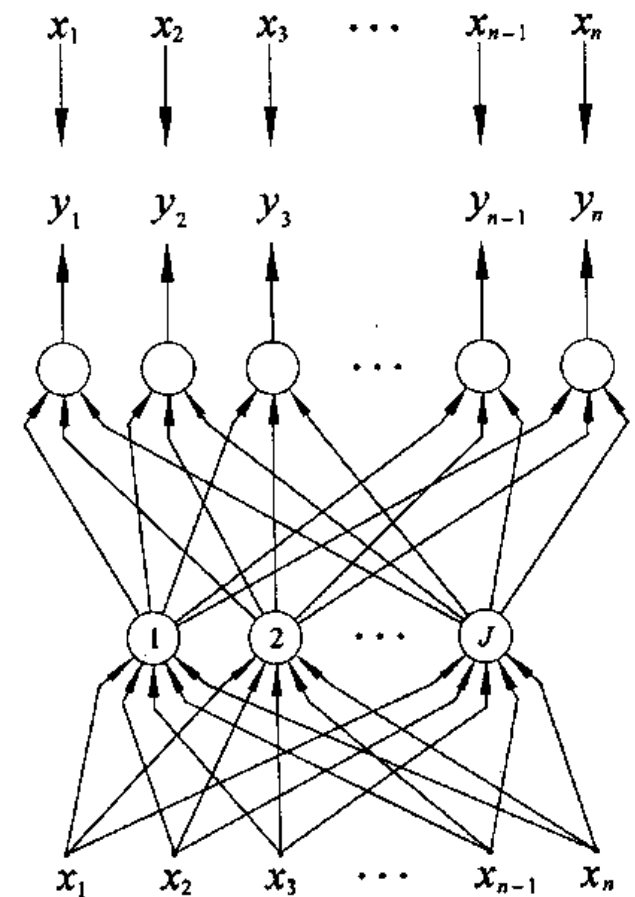
$$E = \sum_{p=1}^P E^p(\mathbf{x}) = \sum_{p=1}^P \left(\alpha^p - N(\mathbf{x}^p) \right)^2$$

- 1 hidden layer
- 10 hidden units
- training set: 9000 pairs
- test set: 3000 pairs
- on-line training

- Results:
 - Physical characterisation: RMS = 53.6 %
 - MLP learning: RMS = 34.9 %

Application: image compression 1

- lossy compression scheme (example)
- original image split into 8x8 pixel regions
- each 64-features vector sent in auto-associative MLP
- hidden layer: 16 neurons
- compression ratio: 16/64



From: M. Hassoun, Fundamentals of artificial neural networks, MIT Press, 1995

Application: image compression 1

- result on training image
 - (a) original
 - (b) compressed
 - (c) compressed + quantized (1.5 bit/pixel)
 - (d) compressed + quantized (1 bit/pixel)



(a)



(b)



(c)




(d)


From: M. Hassoun, Fundamentals of artificial neural networks, MIT Press, 1995

Application: image compression 1


- result on testimage
 - (a) original
 - (b) compressed
 - (c) compressed + quantized (1.5 bit/pixel)
 - (d) compressed + quantized (1 bit/pixel)



(a)



(b)



(c)




(d)

From: M. Hassoun, Fundamentals of artificial neural networks, MIT Press, 1995

Application: image compression 2


- lossless compression scheme (example)
- *Laplacian pyramids*



- If output and error are transmitted: lossless compression !
(interpolator function also needed of course)

Application: image compression 2

- Decimator: (half-band filter) + downsampler ← anti-aliasing




- Interpolator: upsampler + (half-band filter) ← smoothing

Application: image compression 2

- Traditional Laplace pyramids: decimator and interpolator are linear
- MLP: interpolator is non-linear
 - better reconstruction
 - smaller error
 - better compression ratio for error with entropic coder
- entropic coder: better ratio if more zeroes

Application: image compression 2




From "Laplacian pyramid with multilayer perceptrons interpolators", B. Simon, B. Macq, M. Verleysen, in ESANN'93 (Bruges, Belgium) proceedings, D-Facto publications, 1993

Application: image compression 2


linear



MLP

reconstructed
images


differences



From "Laplacian pyramid with multilayer perceptrons interpolators", B. Simon, B. Macq, M. Verleysen, in ESANN'93 (Bruges, Belgium) proceedings, D-Facto publications, 1993

Application: time series forecasting

- Time series:




- Applications:

- finance (stock market index, exchange rates, etc.)
- electricity / gas consumption
- etc.

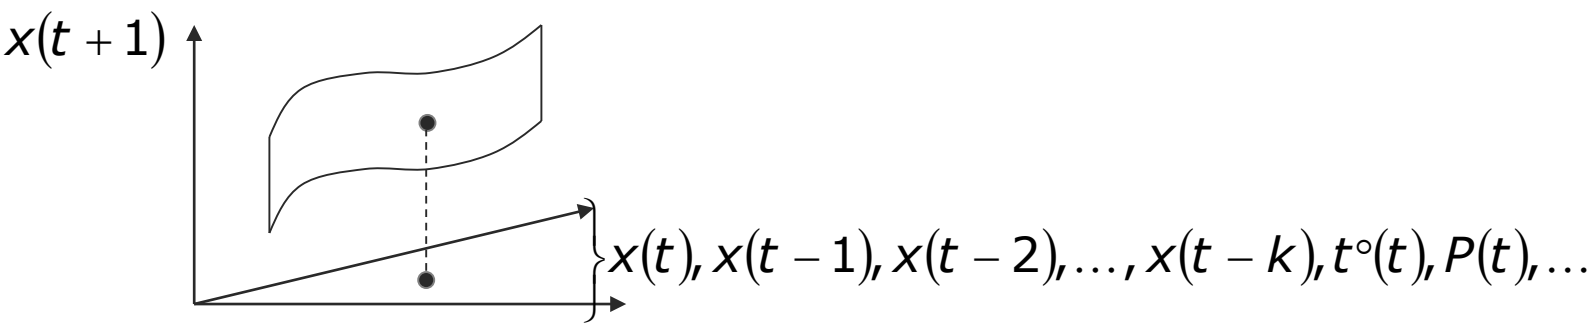
Application: time series forecasting

- Time series forecasting = function approximation



$$x(t + 1) = f(\underbrace{x(t), x(t - 1), x(t - 2), \dots, x(t - k)}_{\text{past values}}, \underbrace{t^\circ(t), P(t), \dots}_{\text{inputs (exogenous variables)}})$$


- New point in *inside* the surface



- So it is really interpolation, not extrapolation!

Application: time series forecasting

- Long-term forecasting



- To predict $x(t+2)$ we can use
 - the true value
 - the estimated value of $x(t+1)$