

ELEC2870 - Machine learning: regression and dimensionality reduction

Nonlinear regression with Radial-Basis Function Networks

Michel Verleysen

Machine Learning Group

Université catholique de Louvain

Louvain-la-Neuve, Belgium

michel.verleysen@uclouvain.be

Outline

- Origin: Cover's theorem
- Interpolation problem
- Regularization theory
- Generalized RBFN
 - Universal approximation
 - RBFN = kernel regression
 - Comparison with MLP
- Learning
 - Centers
 - Widths
 - Multiplying factors
 - Other forms

Origin: Covers' theorem

- Covers' theorem on separability of patterns (1965)
- x^1, x^2, \dots, x^P assigned to two classes C^1, C^2
- ϕ -separability:

$$\exists w \mid \begin{cases} w^T \phi(x) > 0 & x \in C^1 \\ w^T \phi(x) < 0 & x \in C^2 \end{cases}$$

- Cover's theorem:
 - if functions $\phi_i(x)$ are nonlinear
 - if number of functions $\phi_i(x) > \text{dimension input space}$
 - then probability of separability closer to 1
- Quite a « natural » theorem:
 - imagine that you take many (really, a lot...) of random transformations $\phi_i(x)$. They form a (very) high-dimensional space, and with some chance, some features in that space will be linearly separable!

Outline

- Origin: Cover's theorem
- Interpolation problem
- Regularization theory
- Generalized RBFN
 - Universal approximation
 - RBFN = kernel regression
 - Comparison with MLP
- Learning
 - Centers
 - Widths
 - Multiplying factors
 - Other forms

Interpolation problem

- Given points (x^p, t^p) , $x^p \in \mathbb{R}^D$, $t^p \in \mathbb{R}$, $1 \leq p \leq P$:

- Find $F : \mathbb{R}^D \rightarrow \mathbb{R}$ that satisfies

$$F(x^p) = t^p, \quad p = 1 \dots P$$

(so we don't care for regularization, for now)

- RBF technique (Powell, 1988):

$$F(x) = \sum_{p=1}^P w_p \varphi(\|x - x^p\|)$$

- $\varphi(\|x - x^p\|)$ are arbitrary non-linear functions (RBF)
- as many functions as data points
- centers fixed at known points x^p

Interpolation problem

$$F(x^p) = t^p \qquad F(x) = \sum_{p=1}^P w_k \varphi(\|x - x^p\|)$$

$$\begin{bmatrix} \varphi_{11} & \varphi_{12} & \cdots & \varphi_{1P} \\ \varphi_{21} & \varphi_{22} & \cdots & \varphi_{2P} \\ \vdots & \vdots & \ddots & \vdots \\ \varphi_{P1} & \varphi_{P2} & \cdots & \varphi_{PP} \end{bmatrix} \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_P \end{bmatrix} = \begin{bmatrix} t^1 \\ t^2 \\ \vdots \\ t^P \end{bmatrix}$$

where

$$\varphi_{kl} = \varphi(\|x^k - x^l\|)$$

- Into matrix form: $\Phi w = t \rightarrow w = \Phi^{-1}t$
- Vital question: is Φ non-singular ?

Michelli's theorem

- If points \mathbf{x}^k are distinct, Φ is non-singular (regardless of the dimension of the input space)
- Valid for a large class of RBF functions:


$\varphi(\mathbf{x}, \mathbf{c}) = \sqrt{\ \mathbf{x} - \mathbf{c}\ ^2 + l^2}$	$(l > 0)$	non-localized function localized functions
$\varphi(\mathbf{x}, \mathbf{c}) = \frac{1}{\sqrt{\ \mathbf{x} - \mathbf{c}\ ^2 + l^2}}$		
$\varphi(\mathbf{x}, \mathbf{c}) = \exp\left(-\frac{\ \mathbf{x} - \mathbf{c}\ ^2}{2\sigma^2}\right)$	$(\sigma > 0)$	

Outline

- Origin: Cover's theorem
- Interpolation problem
- Regularization theory
- Generalized RBFN
 - Universal approximation
 - RBFN = kernel regression
 - Comparison with MLP
- Learning
 - Centers
 - Widths
 - Multiplying factors
 - Other forms

But...

- Remember: learning is all in-posed problem
- Nobody cares (or should...) about learning; the challenge is generalization!



- Error criterion:
$$E(F) = \underbrace{\frac{1}{2P} \sum_{p=1}^P (t^p - F(x^p))^2}_{\text{MSE}} + \lambda \underbrace{\frac{1}{2} C(w)}_{\text{regularization}}$$

Solution to the regularization problem

- Poggio & Girosi (1990):
 - if $C(w)$ is a (problem-dependent) linear differential operator, the solution to

$$E(F) = \frac{1}{2P} \sum_{p=1}^P (t^p - F(x^p))^2 + \lambda \frac{1}{2} C(w)$$

is of the following form:

$$F(x) = \sum_{p=1}^P w_p G(x, x^p)$$

where $G()$ is a Green's function,

$$w = (G + \lambda I)^{-1} t$$

$$G_{kl} = G(x^k, x^l)$$

Interpolation - regularization

- Interpolation

$$F(x) = \sum_{p=1}^P w_p \varphi(\|x - x^p\|)$$

$$w = \Phi^{-1}t$$

- Exact interpolator

- Possible RBF:

$$\varphi(\|x, x^p\|) = \exp\left(-\frac{\|x - x^p\|^2}{2\sigma^2}\right)$$

- Regularization

$$F(x) = \sum_{p=1}^P w_p G(x, x^p)$$

$$w = (G + \lambda I)^{-1}t$$

- Exact interpolator

- Equal to the interpolation solution iff $\lambda=0$

- Example of Green function:

$$G(x, x^p) = \exp\left(-\frac{\|x - x^p\|^2}{2\sigma^2}\right)$$

One RBF / Green's function for each learning pattern!

Outline

- Origin: Cover's theorem
- Interpolation problem
- Regularization theory
- Generalized RBFN
 - Universal approximation
 - RBFN = kernel regression
 - Comparison with MLP
- Learning
 - Centers
 - Widths
 - Multiplying factors
 - Other forms

Generalized RBFN (GRBFN – RBFN)

- As many radial functions as learning patterns:
 - computationally (too) intensive
(inversion of $P \times P$ matrix grows with P^3)
 - ill-conditioned matrix
 - regularization not easy (problem-specific)
- *Generalized* RBFN approach!

$$F(x) = \sum_{i=1}^K w_i \varphi(\|x - c_i\|)$$

Typically:

- $K \ll P$


- $\varphi(\|x - c_i\|) = \exp\left(-\frac{\|x - c_i\|^2}{2\sigma_i^2}\right)$

Parameters:
 c_i, σ_i, w_i

Radial-Basis Function Networks (RBFN)

$$F(x) = \sum_{i=1}^K w_i \varphi(\|x - c_i\|)$$

$$\varphi(\|x - c_i\|) = \exp\left(-\frac{\|x - c_i\|^2}{2\sigma_i^2}\right)$$



- Possibilities:
 - several outputs (common hidden layer)
 - bias (recommended) (see extensions)

RBFN: universal approximation

- Park & Sandberg 1991:

- For any continuous input-output mapping function $f(\mathbf{x})$

$$\exists F(\mathbf{x}) = \sum_{i=1}^K w_i \varphi(\|\mathbf{x} - \mathbf{c}_i\|) \mid L_p(f(\mathbf{x}), F(\mathbf{x})) < \varepsilon \quad (\varepsilon > 0, p \in [1, \infty])$$

- The theorem is stronger (radial symmetry not needed)
- K not specified
- Provides a theoretical basis for *practical* RBFN!

RBFN and kernel regression

- A digression about (probabilistic) kernel regression:
- non-linear regression model

$$t^p = f(\mathbf{x}^p) + \varepsilon^p = y^p + \varepsilon^p, \quad 1 \leq p \leq P$$

- estimation of $f(\mathbf{x})$: average of t around \mathbf{x} . More precisely:

$$\begin{aligned} f(x) &= E[y|x] \\ &= \int_{-\infty}^{\infty} y f_Y(y|x) dy \\ &= \frac{\int_{-\infty}^{\infty} y f_{X,Y}(x,y) dy}{f_X(x)} \end{aligned}$$

- Need for estimates of $f_{X,Y}(x,y)$ and $f_X(x)$
→ Parzen-Rosenblatt density estimator

Parzen-Rosenblatt density estimator

$$\hat{f}_x(x) = \frac{1}{Ph^d} \sum_{p=1}^P K\left(\frac{x - x^p}{h}\right)$$

with $K()$ continuous, bounded, symmetric about the origin, with maximum value at 0, and with unit integral, is consistent (asymptotically unbiased).


- Estimation of $f_{X,Y}(x,y)$

$$\hat{f}_{X,Y}(x,y) = \frac{1}{Ph^{d+1}} \sum_{p=1}^P K\left(\frac{x - x^p}{h}\right) K\left(\frac{y - y^p}{h}\right)$$

RBFN and kernel regression

$$\hat{f}(x) = \frac{\int_{-\infty}^{\infty} y \hat{f}_{X,Y}(x,y) dy}{\hat{f}_X(x)}$$

$$= \frac{\sum_{p=1}^P y^p K\left(\frac{x - x^p}{h}\right)}{\sum_{p=1}^P K\left(\frac{x - x^p}{h}\right)}$$

$$f(x) = \frac{\int_{-\infty}^{\infty} y f_{X,Y}(x,y) dy}{f_X(x)}$$


- Weighted average of y^i
- called Nadaraya-Watson estimator (1964)
- equivalent to *Normalized RBFN* in the unregularized context

RBFN - MLP

- RBFN

- single hidden layer
- non-linear hidden layer
linear output layer
- argument of hidden units:
Euclidean norm
- universal approximation
property
- local approximators
- splitted learning

- MLP

- single or multiple hidden layers
- non-linear hidden layer
linear or non-linear output layer
- argument of hidden units:
scalar product
- universal approximation
property
- global approximators
- global learning

Outline

- Origin: Cover's theorem
- Interpolation problem
- Regularization theory
- Generalized RBFN
 - Universal approximation
 - RBFN = kernel regression
 - Comparison with MLP
- Learning
 - Centers
 - Widths
 - Multiplying factors
 - Other forms

RBFN learning

- This doesn't tell us how to *learn* in a RBFN

$$F(x) = \sum_{i=1}^K w_i \varphi(\|x - c_i\|) \qquad \varphi(\|x - c_i\|) = \exp\left(-\frac{\|x - c_i\|^2}{2\sigma_i^2}\right)$$


- Remember: 3 set of parameters c_i, σ_i, w_i
- Traditional learning strategy: splitted computation
 1. centers c_i
 2. widths σ_i
 3. weights w_i

RBFN step 1: centers

- Idea: density of centers c_i must follow the density of learning points x^k
→ vector quantization
 - selected at random (in learning set)
 - competitive learning
 - frequency-sensitive learning
 - Kohonen maps
 - ...
- This phase only uses the x^k information, not the t^k (it is *unsupervised*)

RBFN step 2: widths

- Universal approximation property holds with identical widths
- In practice (limited learning set): variable widths σ_i
- Idea: RBFN use *local* clusters



- choose σ_i *according* to standard deviation of clusters
- *According* means high stdv of clusters, high radius, but *doesn't* mean equality (see further for details and ideas)

RBFN step 3: weights

$$F(x) = \sum_{i=1}^K w_i \varphi(\|x - c_i\|) \qquad \varphi(\|x - c_i\|) = \exp\left(-\frac{\|x - c_i\|^2}{2\sigma_i^2}\right)$$

at x^p : constants !

- Problem becomes **linear** !
- Solution of least square criterion leads to

$$E(F) = \frac{1}{2P} \sum_{p=1}^P (t^p - F(x^p))^2$$

$$w = \Phi^+ t = (\Phi^T \Phi)^{-1} \Phi^T t$$

where

$$\Phi \equiv \varphi_{ki} = \varphi(\|x^k - c_i\|)$$

- In practise: use SVD !

RBFN step 4: gradient descent

- After steps 1 to 3: all parameters are computed
 - But only the weights w are optimized with respect to the error criterion
 - Centers and widths are *reasonable choices*

$$F(x) = \sum_{i=1}^K w_i \exp\left(-\frac{\|x - c_i\|^2}{2\sigma_i^2}\right)$$

Diagram illustrating the components of the RBFN equation:

- 1**: The entire equation $F(x) = \sum_{i=1}^K w_i \exp\left(-\frac{\|x - c_i\|^2}{2\sigma_i^2}\right)$ is supervised.
- 2**: The Gaussian kernel $\exp\left(-\frac{\|x - c_i\|^2}{2\sigma_i^2}\right)$ is unsupervised.
- 3**: The weight w_i is supervised.

- Optional step: possibility of gradient descent on *all* parameters
- Some improvement, but
 - learning speed
 - local minima
 - risk of non-local basis functions
 - etc.

More elaborated models

- Add constant and linear terms

$$F(x) = \sum_{i=1}^K w_i \exp\left(-\frac{\|x - c_i\|^2}{2\sigma_i^2}\right) + \sum_{i=1}^D w'_i x_i + w'_0$$

good idea (very difficult to approximate a constant with kernels...)

- Use normalized RBFN

$$F(x) = \sum_{i=1}^K w_i \frac{\exp\left(-\frac{\|x - c_i\|^2}{2\sigma_i^2}\right)}{\sum_{j=1}^K \exp\left(-\frac{\|x - c_j\|^2}{2\sigma_j^2}\right)}$$

basis functions are bounded [0,1] → can be interpreted as probability values (classification)

Back to the widths...

- choose σ_i according to standard deviation of clusters
- In the literature:
 - $\sigma = d_{\max} / \sqrt{2K}$ where d_{\max} = maximum distance between centroids [1]
 - $\sigma_i = \frac{1}{q} \sqrt{\sum_{j=1}^q \|c_i - c_j\|^2}$ where index j scans the q nearest centroids to c_i [2]
 - $\sigma_i = r \min_j (\|c_i - c_j\|)$ where r is an overlap constant [3]
 -





[1] S. Haykin, *"Neural Networks a Comprehensive Foundation"*, Prentice-Hall Inc, second edition, 1999.

[2] J. Moody and C. J. Darken, *"Fast learning in networks of locally-tuned processing units"*, Neural Computation 1, pp. 281-294, 1989.

[3] A. Saha and J. D. Keeler, *"Algorithms for Better Representation and Faster Learning in Radial Basis Function Networks"*, Advances in Neural Information Processing Systems 2, Edited by David S. Touretzky, pp. 482-489, 1989.

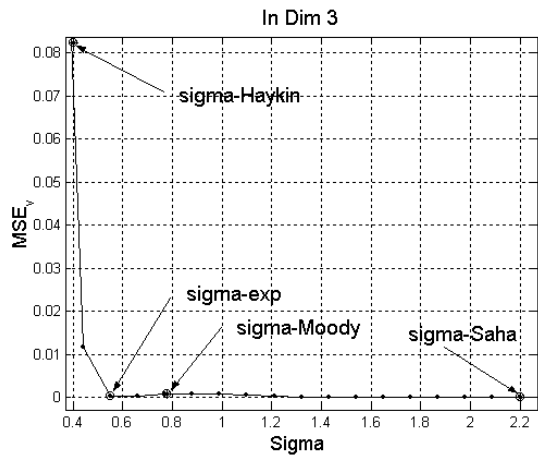
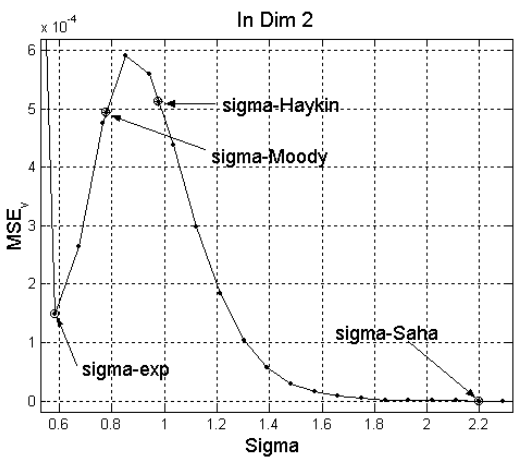
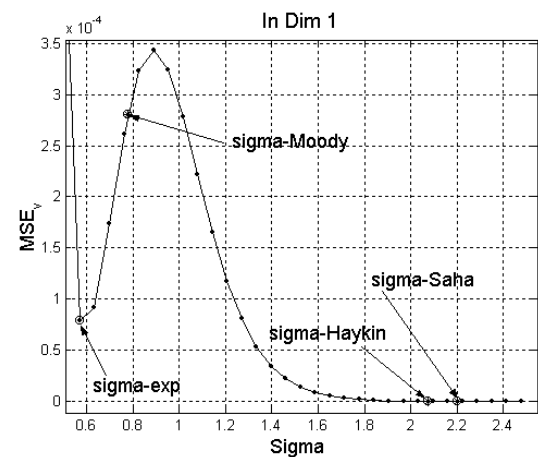
Basic example

- Approximation of $f(\mathbf{x}) = 1$ with a d -dimensional RBFN
- In theory: identical w_i
- Experimentally: side effects
→ only middle taken into account

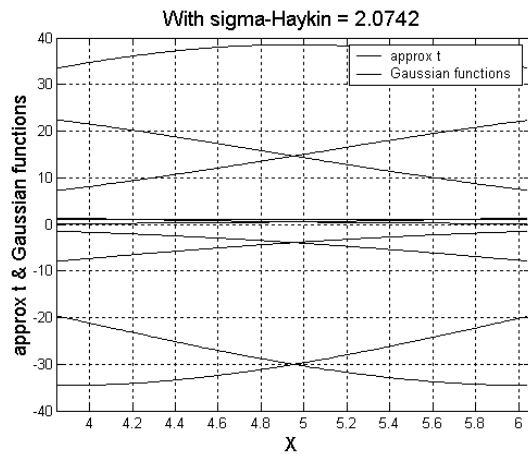

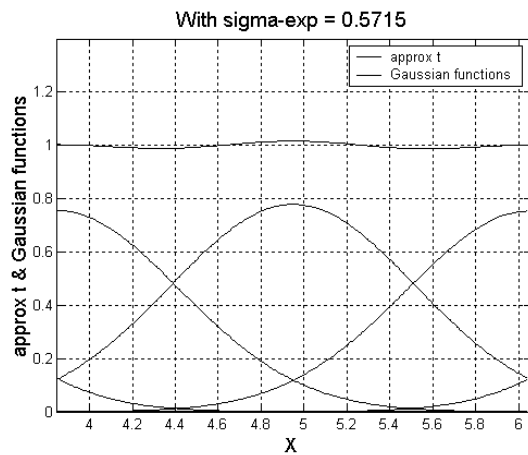


Error versus width

Basic example: error versus space dimension

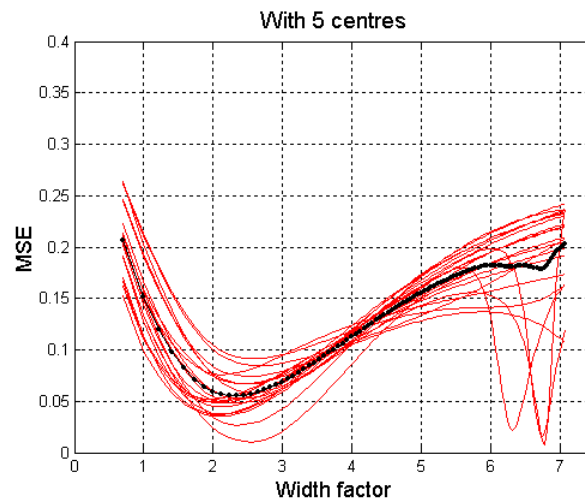


Basic example: local decomposition?



Multiple local minima in error curve

- Choose the first minimum to preserve the locality of clusters



- The first local minimum is usually less sensitive to variability (sparsity is good!)

Some concluding comments

- RBFN: easy learning (compared to MLP)
 - in a cross-validation scheme: important!
- Many RBFN models
- Even more RBFN learning schemes...
- Results not very sensitive to unsupervised part of learning (c_i, σ_i)
- Open work for a priori (problem-dependent) choice of widths σ_i

Sources and references

- Most of the basic concepts developed in these slides come from the excellent book:
 - Neural networks – a comprehensive foundation, S. Haykin, Macmillan College Publishing Company, 1994.
- Some supplementary comments come from the tutorial on RBF:
 - An overview of Radial Basis Function Networks, J. Ghosh & A. Nag, in: Radial Basis Function Networks 2, R.J. Howlett & L.C. Jain eds., Physica-Verlag, 2001.